Loadsmart ABM OS

A ready-to-run kit for Loadsmart's AE/SDR Account Based Selling pods: one daily pipeline from signal to sequenced outreach, three products (Opendock, ShipperGuide, Loadsmart AI) kept to their own buyers, approval gates on every send and CRM write, and a dashboard that shows where the program leaks.

Built with Enably (SalesGency) by Patrick Diamitani, November 2026. The full playbook (ICP, personas, messaging, cadences, KPIs, rhythm) is the companion doc.

What's inside

Folder	What it is
`skills/`	11 skills, one per pipeline step, plus an orchestrator
`skills-zips/`	The same skills zipped one per file, ready to upload to claude.ai (Settings > Capabilities > Skills)
`config/`	ICP, persona map, pods and caps, sending and dialer setup, daily schedule, Salesforce fields
`agent/`	Pod Copilot system prompt (Claude Project or any agent) and `CLAUDE.md` for Claude Code
`dashboard/index.html`	Pod dashboard. Opens in any browser; import the daily CSV; nothing leaves the browser
`templates/`	Dashboard CSV template, disqualification reason codes

The pipeline

When	Skill	Does
Week 1, once	`la-crm-audit`	Inventories and ICP-scores the existing CRM so outreach starts from what Loadsmart owns
Daily 06:00	`la-signal`	Dated, sourced why-now signals mapped to a product
Weekly Mon	`la-abm-tag`	20 accounts per pod with written scores (RevOps approves)
Daily 07:00	`la-persona-enrich`	Verified contacts per product persona, gaps only
Daily 07:30	`la-research`	Sourced brief per contact
Daily 08:00	`la-email`	Body-only emails 1-3, QA-checked
Daily 08:30	`la-sequence`	Pod approves batch; Gong flow with email, dialer + voicemail drop, LinkedIn
Daily 17:00	`la-daily-report`	Funnel, leak flags, digest
Weekly	`la-list-qa`, `la-pod-feedback`	Quality scorecard; rep feedback turned into shipped changes

Setup, in order (no coding)

1. **RevOps:** create the Salesforce fields in `config/salesforce-fields.md`.
2. **Fill the placeholders** in `config/` (anything in `[brackets]` or marked ASSUMPTION). Marketing fills `approved_proof_points` in `personas.yaml`.
3. **Sending:** buy secondary sending domains, create 5 inboxes per rep, set SPF/DKIM/DMARC, start warm-up on day 1 (2-3 weeks before any cold send). See `config/sending.yaml`.
4. **Orchestrator:** in n8n (or Make.com), create one scheduled workflow per row of `config/schedule.yaml`. Rule-based steps are workflow nodes; Claude is one node for research, copy and scoring rationale. Store keys in the tool's credential store, never in a config file.
5. **Skills:** upload the zips in `skills-zips/` to Claude, or point Claude Code at this folder (`agent/CLAUDE.md`).

6. **Copilot:** create a Claude Project, paste `agent/pod-copilot-system-prompt.md` into its instructions, add `config/` and `skills/` as knowledge.
7. **Run `la-crm-audit` first**, then a dry run of the daily pipeline with 3 accounts and nothing sent.
8. **Legal review** of cold email footer, voicemail drop and any Canadian contacts before going live.
9. **Dashboard:** open `dashboard/index.html`, click Import CSV and load the file `la-daily-report` produces.

Ground rules

- Nothing sends, enrolls or writes to Salesforce without a human yes for that batch.
- Nothing in copy is invented. Missing facts stay as `[placeholders]`; website claims are used only once Marketing approves them.
- One product and one persona per email. Cross-product touches per account are capped.
- Thresholds in the skills are starting points; replace them with Loadsmart's 30-day baseline.